Applied AI / Machine Learning Engineer

Full-time | Hybrid / Remote | Sample role specification (October 2026)

JOB SUMMARY

We are looking for an Applied AI / Machine Learning Engineer who can build intelligent applications, machine-learning workflows, and LLM-powered products. The role combines hands-on model development, practical software integration, and evaluation of AI solutions.

REQUIRED EXPERIENCE

Minimum 3 years of professional experience developing or delivering applied AI, machine-learning, deep-learning, or LLM-powered software. Academic coursework, personal learning, and university projects may demonstrate individual technical skills but do not independently satisfy the professional-duration requirement.

REQUIRED SKILLS

- Python
- Machine Learning
- Deep Learning
- TensorFlow or PyTorch (at least one)
- Natural Language Processing (NLP)
- LLM integrations
- LangChain
- Model evaluation
- AI application development

PREFERRED SKILLS

- LangGraph
- Medical imaging
- Time-series forecasting
- Google Cloud
- Node.js
- Firebase
- Experience building production-oriented AI products

CORE RESPONSIBILITIES

- Develop Python-based AI applications and reproducible ML/deep-learning training or inference workflows.
- Integrate large language models into useful products, services, or internal tools.
- Use NLP and retrieval or orchestration frameworks where relevant to the use case.
- Evaluate solutions with appropriate validation data, baseline comparisons, error analysis, and documented metrics.
- Collaborate with engineers to expose reliable APIs, monitor deployments, and improve model behavior.

HOW TO INTERPRET REQUIREMENTS

- The TensorFlow/PyTorch condition is OR: evidence of either framework can satisfy that named-tool requirement. All other listed required skills are evaluated individually.
- Equivalent functionality demonstrated with different tools may be presented as potential transferable capability; it is not automatically proof of experience with a specifically named technology.
- The 3-year requirement concerns the relevant professional work itself: unrelated tenure should not be counted as AI/ML engineering experience.
- Preferred skills are additional strengths, not mandatory gates.